

IMPORTING LIBRARIES

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

IMPORTING FILES FOR ANALYSIS

```
In [2]: trans=pd.read_csv("Retail_Data_Transactions.csv")
trans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 125000 entries, 0 to 124999
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   customer_id     125000 non-null  object
1   trans_date      125000 non-null  object
2   tran_amount     125000 non-null  int64
dtypes: int64(1), object(2)
memory usage: 2.9+ MB
```

```
In [3]: resp=pd.read_csv("Retail_Data_Response.csv")
resp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6884 entries, 0 to 6883
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   customer_id     6884 non-null  object
1   response        6884 non-null  int64
dtypes: int64(1), object(1)
memory usage: 107.7+ KB
```

MERGING BOTH THE CSV(s)

```
In [4]: df=trans.merge(resp,on='customer_id',how='left')
df
```

```
Out[4]:
```

	customer_id	trans_date	tran_amount	response
0	CS5295	2013-02-11	35	1.0
1	CS4768	2015-03-15	39	1.0
2	CS2122	2013-02-26	52	0.0
3	CS1217	2011-11-16	99	0.0
4	CS1850	2013-11-20	78	0.0
...
124995	CS8433	2011-06-26	64	0.0
124996	CS7232	2014-08-19	38	0.0
124997	CS8731	2014-11-28	42	0.0
124998	CS8133	2013-12-14	13	0.0
124999	CS7996	2014-12-13	36	0.0

125000 rows × 4 columns

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 125000 entries, 0 to 124999
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   customer_id     125000 non-null object
1   trans_date      125000 non-null object
2   tran_amount     125000 non-null int64
3   response        124969 non-null float64
dtypes: float64(1), int64(1), object(2)
memory usage: 3.8+ MB
```

HANDELING MISSING VALUE(s)

```
In [6]: df.isnull().sum()
```

```
Out[6]: customer_id    0
trans_date    0
tran_amount    0
response      31
dtype: int64
```

```
In [7]: df=df.dropna()
df.reset_index()
```

Out[7]:

	index	customer_id	trans_date	tran_amount	response	
	0	0	CS5295	2013-02-11	35	1.0
	1	1	CS4768	2015-03-15	39	1.0
	2	2	CS2122	2013-02-26	52	0.0
	3	3	CS1217	2011-11-16	99	0.0
	4	4	CS1850	2013-11-20	78	0.0

124964	124995	CS8433	2011-06-26	64	0.0	
124965	124996	CS7232	2014-08-19	38	0.0	
124966	124997	CS8731	2014-11-28	42	0.0	
124967	124998	CS8133	2013-12-14	13	0.0	
124968	124999	CS7996	2014-12-13	36	0.0	

124969 rows × 5 columns

MANAGING DATA-TYPE(s)

In [8]:

```
df['trans_date']=df['trans_date'].astype(dtype='datetime64[ns]')
df['response']=df['response'].astype(dtype='int64')
df['tran_amount']=df['tran_amount'].astype(dtype='float64')
df
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\3299215634.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['trans_date']=df['trans_date'].astype(dtype='datetime64[ns]')
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\3299215634.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['response']=df['response'].astype(dtype='int64')
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\3299215634.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['tran_amount']=df['tran_amount'].astype(dtype='float64')
```

Out[8]:

	customer_id	trans_date	tran_amount	response
0	CS5295	2013-02-11	35.0	1
1	CS4768	2015-03-15	39.0	1
2	CS2122	2013-02-26	52.0	0
3	CS1217	2011-11-16	99.0	0
4	CS1850	2013-11-20	78.0	0
...
124995	CS8433	2011-06-26	64.0	0
124996	CS7232	2014-08-19	38.0	0
124997	CS8731	2014-11-28	42.0	0
124998	CS8133	2013-12-14	13.0	0
124999	CS7996	2014-12-13	36.0	0

124969 rows × 4 columns

SEARCHING FOR OUTLIER(s)

In [9]:

```
# Z-Score Method
z1_scores=(df['tran_amount']-df['tran_amount'].mean())/df['tran_amount'].std()
z1_threshold=3
outliers_1 = df[abs(z1_scores)>z1_threshold]
if outliers_1.empty:
    print("NO OUTLIERS PRESENT IN '{}' COLUMN".format('tran_amount'))
else:
    print("THE OUTLIERS PRESENT IN COLUMN '{}' ARE AS FOLLOWS \n".format('tran_a
    print(outliers_1)

z2_scores=(df['response']-df['response'].mean())/df['response'].std()
z2_threshold=3
outliers_2 = df[abs(z2_scores)>z2_threshold]
if outliers_2.empty:
    print("NO OUTLIERS PRESENT IN '{}' COLUMN".format('response'))
else:
    print("THE OUTLIERS PRESENT IN COLUMN '{}' ARE AS FOLLOWS \n".format('respon
    print(outliers_2)
```

NO OUTLIERS PRESENT IN 'tran_amount' COLUMN

NO OUTLIERS PRESENT IN 'response' COLUMN

In [10]:

```
# IQR(Inter Quartile Range) Method
q1=df['tran_amount'].quantile(0.25)
q3=df['tran_amount'].quantile(0.75)
iqr_1=q3-q1
lower_thres1=q1-1.5*iqr_1
upper_thres1=q3+1.5*iqr_1
outliers_1=df[(df['tran_amount']<lower_thres1)|(df['tran_amount']>upper_thres1)]
if outliers_1.empty:
    print("NO OUTLIERS PRESENT IN '{}' COLUMN".format('tran_amount'))
else:
    print("THE OUTLIERS PRESENT IN COLUMN '{}' ARE AS FOLLOWS: \n".format('tran_
```

```

print(outliers_1)

print("\n\n")

q1=df['response'].quantile(0.25)
q3=df['response'].quantile(0.75)
iqr_2=q3-q1
lower_thres2=q1-1.5*iqr_2
upper_thres2=q3+1.5*iqr_2
outliers_2=df[(df['response']<lower_thres2)|(df['response']>upper_thres2)]
if outliers_2.empty:
    print("NO OUTLIERS PRESENT IN '{}' COLUMN".format('response'))
else:
    print("THE OUTLIERS PRESENT IN COLUMN '{}' ARE AS FOLLOWS: \n".format('response'))
    print(outliers_2)

```

NO OUTLIERS PRESENT IN 'tran_amount' COLUMN

THE OUTLIERS PRESENT IN COLUMN 'response' ARE AS FOLLOWS:

	customer_id	trans_date	tran_amount	response
0	CS5295	2013-02-11	35.0	1
1	CS4768	2015-03-15	39.0	1
9	CS3802	2013-08-20	75.0	1
25	CS2748	2013-03-23	37.0	1
35	CS5514	2013-01-09	36.0	1
...
124922	CS7182	2011-07-07	45.0	1
124923	CS8845	2013-02-23	15.0	1
124946	CS8310	2014-04-23	19.0	1
124986	CS8580	2013-11-15	37.0	1
124988	CS7888	2013-10-19	35.0	1

[13842 rows x 4 columns]


```

In [11]: plt.figure(figsize=(4,4))
sns.boxplot(x=df['tran_amount'],color='orange')
plt.title("Outliers in Transaction Amount")
plt.xlabel("Transaction Amount")
plt.show()


print("\n\n")
plt.figure(figsize=(4,4))
sns.boxplot(x=df['response'],color='red')
plt.title("Outliers in Response Column")
plt.xlabel("Response")
plt.show()

```

Outliers in Transaction Amount



Outliers in Response Column



```
In [12]: df['months']=df['trans_date'].dt.month  
df
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\1941809520.py:1: SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead
 See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 df['months']=df['trans_date'].dt.month

Out[12]:

	customer_id	trans_date	tran_amount	response	months
0	CS5295	2013-02-11	35.0	1	2
1	CS4768	2015-03-15	39.0	1	3
2	CS2122	2013-02-26	52.0	0	2
3	CS1217	2011-11-16	99.0	0	11
4	CS1850	2013-11-20	78.0	0	11
...
124995	CS8433	2011-06-26	64.0	0	6
124996	CS7232	2014-08-19	38.0	0	8
124997	CS8731	2014-11-28	42.0	0	11
124998	CS8133	2013-12-14	13.0	0	12
124999	CS7996	2014-12-13	36.0	0	12

124969 rows × 5 columns

```
In [13]: #Which 3 months have the highest transaction amount?
monthly_sales = df.groupby('months')['tran_amount'].sum().sort_values(ascending=
top3_months = monthly_sales.head(3)
print("Top 3 Months by Total Transaction Amount: \n{}".format(top3_months))
```

Top 3 Months by Total Transaction Amount:

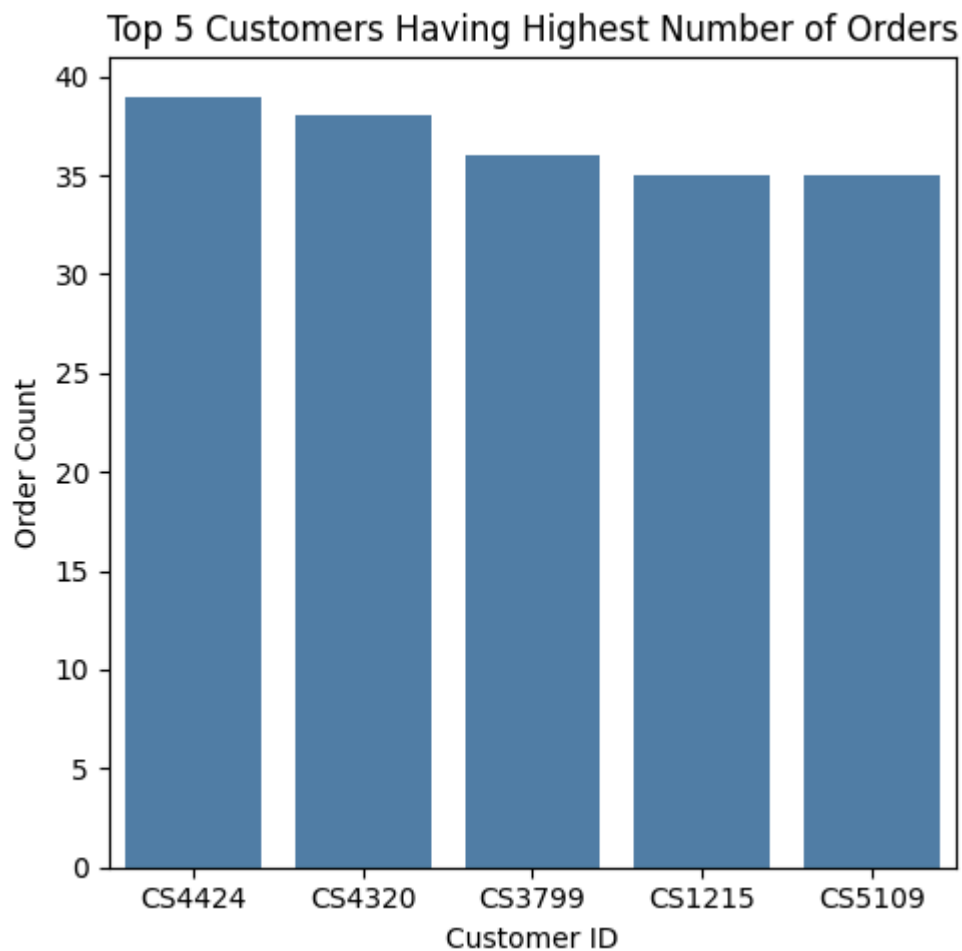
	months	tran_amount
0	8	726775.0
1	10	725058.0
2	1	724089.0

```
In [14]: #Customers having highest number of orders?
customers = df['customer_id'].value_counts().reset_index()
customers.columns = ['customer_id', 'order_count']
top_customers = customers.head()
top_customers
```

Out[14]:

	customer_id	order_count
0	CS4424	39
1	CS4320	38
2	CS3799	36
3	CS1215	35
4	CS5109	35

```
In [15]: plt.figure(figsize=(5,5))
sns.barplot(x='customer_id', y='order_count', data=top_customers,color='steelblue')
plt.title("Top 5 Customers Having Highest Number of Orders")
plt.xlabel("Customer ID")
plt.ylabel("Order Count")
plt.tight_layout()
plt.show()
```




```
In [16]: #Customers having highest value of orders?
customers = df.groupby('customer_id')['tran_amount'].sum().sort_values(ascending=False)
top_customers = customers.head(5)
print("Customers having highest value of orders are: \n{}".format(top_customers))
```


Customers having highest value of orders are:

	customer_id	tran_amount
0	CS4424	2933.0
1	CS4320	2647.0
2	CS5752	2612.0
3	CS4660	2527.0
4	CS3799	2513.0

```
In [17]: plt.figure(figsize=(5, 5))
sns.barplot(x='customer_id', y='tran_amount', data=top_customers,color='orchid')
plt.title("Top 5 Customers with Highest Order Value")
plt.xlabel("Customer ID")
plt.ylabel("Total Transaction Amount")
plt.tight_layout()
plt.show()
```



Advanced analysis

Time Series Analysis

```
In [18]: import matplotlib.dates as mdates
```

```
In [19]: df['month_year']=df['trans_date'].dt.to_period('M')
monthly_sales=df.groupby('month_year')['tran_amount'].sum()
monthly_sales.index=monthly_sales.index.to_timestamp()
print("\n")
plt.figure(figsize=(14,8))
plt.plot(monthly_sales.index,monthly_sales.values, marker='o', linestyle='--')
plt.fill_between(monthly_sales.index, monthly_sales.values, color='skyblue', alp
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
```

```
plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=6))
plt.title('Monthly Sales Trend Over Time', fontsize=15)
plt.xlabel=('Month-Year')
plt.ylabel=('Monthly Sales')
plt.xticks(rotation=45)
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```


C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\3120432034.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['month_year']=df['trans_date'].dt.to_period('M')
```



Cohort Segmentation

```
In [26]: #Recency
recency=df.groupby('customer_id')['trans_date'].max()

#Frequency
frequency=df.groupby('customer_id')['trans_date'].count()

#Monetary
monetary=df.groupby('customer_id')['tran_amount'].sum()

new=pd.DataFrame({'recency':recency,'frequency':frequency,'monetary':monetary})
new.reset_index(inplace=True)
new
```

Out[26]:

	customer_id	recency	frequency	monetary
0	CS1112	2015-01-14	15	1012.0
1	CS1113	2015-02-09	20	1490.0
2	CS1114	2015-02-12	19	1432.0
3	CS1115	2015-03-05	22	1659.0
4	CS1116	2014-08-25	13	857.0
...
6879	CS8996	2014-12-09	13	582.0
6880	CS8997	2014-06-28	14	543.0
6881	CS8998	2014-12-22	13	624.0
6882	CS8999	2014-07-02	12	383.0
6883	CS9000	2015-02-28	13	533.0

6884 rows × 4 columns

In [27]:


```
#Customer Segmentation
def segment_customer(row):
    if (row['recency'].year>=2012) and (row['frequency']>=15) and (row['monetary']>500):
        return 'P0'
    elif (2011<=row['recency'].year<2012) and (10<row['frequency']<15) and (500<row['monetary']<1000):
        return 'P1'
    else:
        return 'P2'
new['Segment']=new.apply(segment_customer,axis=1)
new
```

Out[27]:

	customer_id	recency	frequency	monetary	Segment
0	CS1112	2015-01-14	15	1012.0	P0
1	CS1113	2015-02-09	20	1490.0	P0
2	CS1114	2015-02-12	19	1432.0	P0
3	CS1115	2015-03-05	22	1659.0	P0
4	CS1116	2014-08-25	13	857.0	P2
...
6879	CS8996	2014-12-09	13	582.0	P2
6880	CS8997	2014-06-28	14	543.0	P2
6881	CS8998	2014-12-22	13	624.0	P2
6882	CS8999	2014-07-02	12	383.0	P2
6883	CS9000	2015-02-28	13	533.0	P2

6884 rows × 5 columns

```
In [28]: plt.figure(figsize=(4,4))
sns.countplot(x='Segment', data=new, order=['P0', 'P1', 'P2'], color='lightgreen')
plt.title("Customer Segment Distribution")
plt.xlabel("Segment")
plt.ylabel("Number of Customers")
plt.tight_layout()
plt.show()
```




CHURN ANALYSIS

```
In [31]: #Count the number of churned and active customers
import matplotlib.pyplot as plt

churn_counts = df['response'].value_counts()
churn_counts.plot(kind='bar', color=['salmon', 'skyblue'])

plt.title("Churn vs Active Customers")
plt.xlabel("Response (0 = Churned, 1 = Active)")
plt.ylabel("Number of Customers")
plt.tight_layout()
plt.show()
```



```
In [37]: df['trans_date'] = pd.to_datetime(df['trans_date'])
df['Month'] = df['trans_date'].dt.to_period('M')
monthly_churn = df.groupby(['Month', 'response']).size().unstack(fill_value=0)

monthly_churn.plot(kind='bar', stacked=True, figsize=(14,8), color=['salmon', 'lightblue'])
plt.title("Monthly Churn vs Active Customers")
plt.xlabel("Month")
plt.ylabel("Number of Customers")
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\3253399344.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
df['trans_date'] = pd.to_datetime(df['trans_date'])
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\3253399344.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Month'] = df['trans_date'].dt.to_period('M')
```



```
In [43]: df['trans_date'] = pd.to_datetime(df['trans_date'])
df['Month'] = df['trans_date'].dt.to_period('M')

monthly_counts = df.groupby(['Month', 'response']).size().unstack(fill_value=0)
monthly_counts['Churn Rate (%)'] = monthly_counts[0] / (monthly_counts[0] + mont

monthly_counts['Churn Rate (%)'].plot(kind='line', marker='o', figsize=(14,7), c
plt.title("Monthly Churn Rate (%)")
plt.xlabel("Month")
plt.ylabel("Churn Rate (%)")
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\2235735469.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['trans_date'] = pd.to_datetime(df['trans_date'])
```


C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\2235735469.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

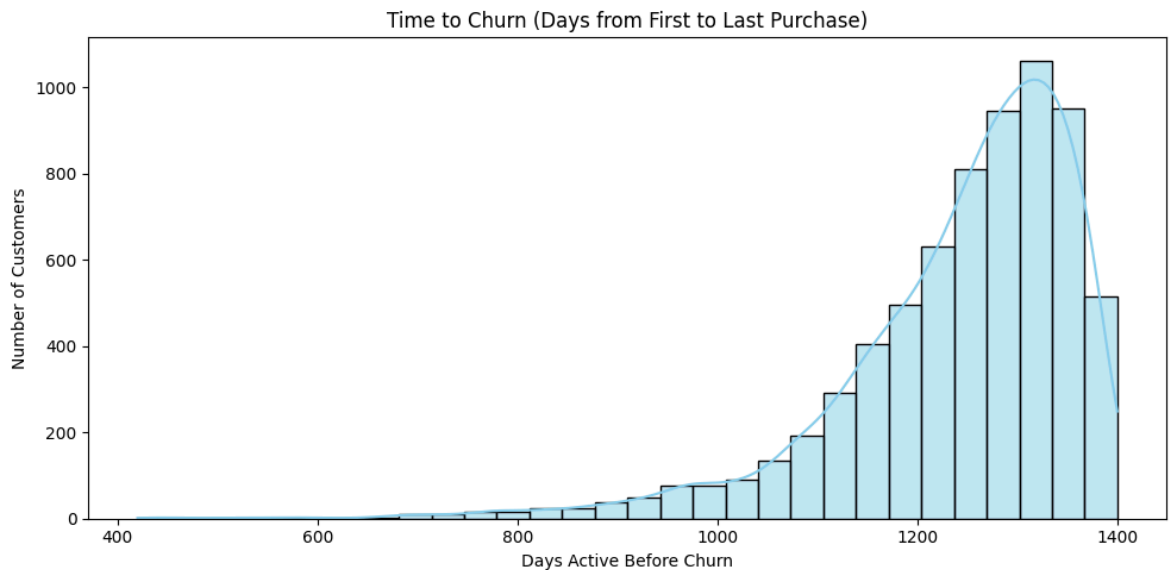
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Month'] = df['trans_date'].dt.to_period('M')
```



```
In [44]: last_purchase = df.groupby('customer_id')['trans_date'].max()
first_purchase = df.groupby('customer_id')['trans_date'].min()
time_to_churn = (last_purchase - first_purchase).dt.days

plt.figure(figsize=(10, 5))
sns.histplot(time_to_churn, bins=30, kde=True, color='skyblue')
plt.title("Time to Churn (Days from First to Last Purchase)")
plt.xlabel("Days Active Before Churn")
plt.ylabel("Number of Customers")
plt.tight_layout()
plt.show()
```



```
In [55]: import matplotlib.pyplot as plt
import seaborn as sns

df['trans_date'] = pd.to_datetime(df['trans_date'])
df['Month'] = df['trans_date'].dt.to_period('M')

churn_table = df.groupby(['Month', 'response']).size().unstack(fill_value=0)
churn_table.columns = ['Churned (0)', 'Active (1)']
churn_table.index = churn_table.index.astype(str)

plt.figure(figsize=(20,10))
sns.heatmap(churn_table.T, annot=True, fmt='d', cmap='YlGnBu', linewidths=0.6, c
```

```
plt.title("Monthly Churn Overview", fontsize=16)
plt.xlabel("Month", fontsize=12)
plt.ylabel("Customer Status", fontsize=12)
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\1818809005.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
df['trans_date'] = pd.to_datetime(df['trans_date'])
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_1596\1818809005.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
df['Month'] = df['trans_date'].dt.to_period('M')
```



ANALYSING THE TOP CUSTOMERS

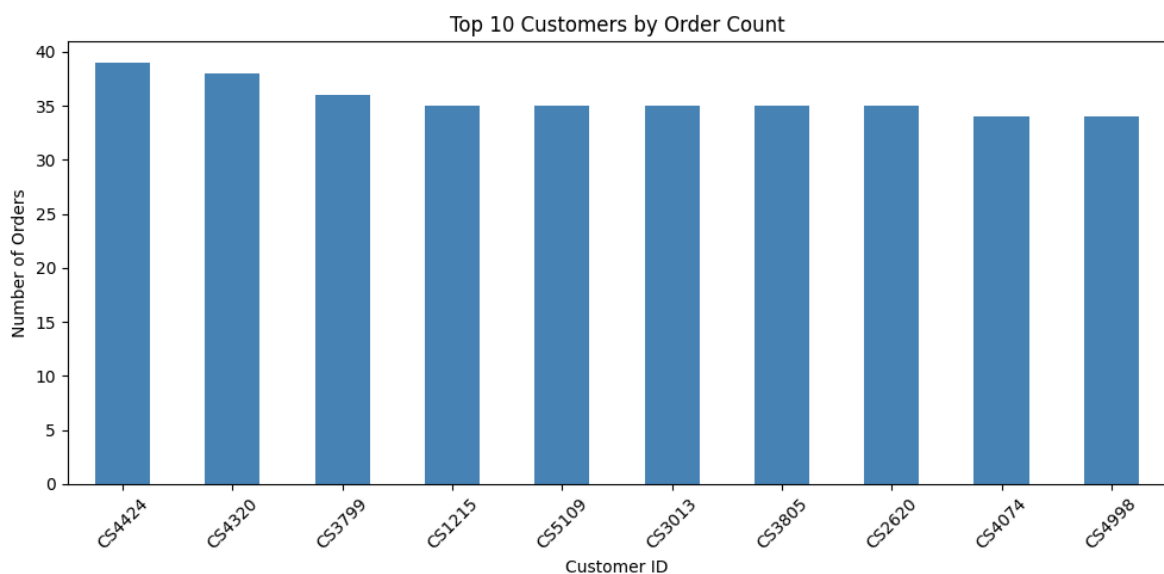
In [56]: `top_spenders = df.groupby('customer_id')['tran_amount'].sum().sort_values(ascending=False)`

```
plt.figure(figsize=(10, 5))
top_spenders.plot(kind='bar', color='mediumseagreen')
plt.title("Top 10 Customers by Total Spend")
plt.xlabel("Customer ID")
plt.ylabel("Total Amount Spent")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

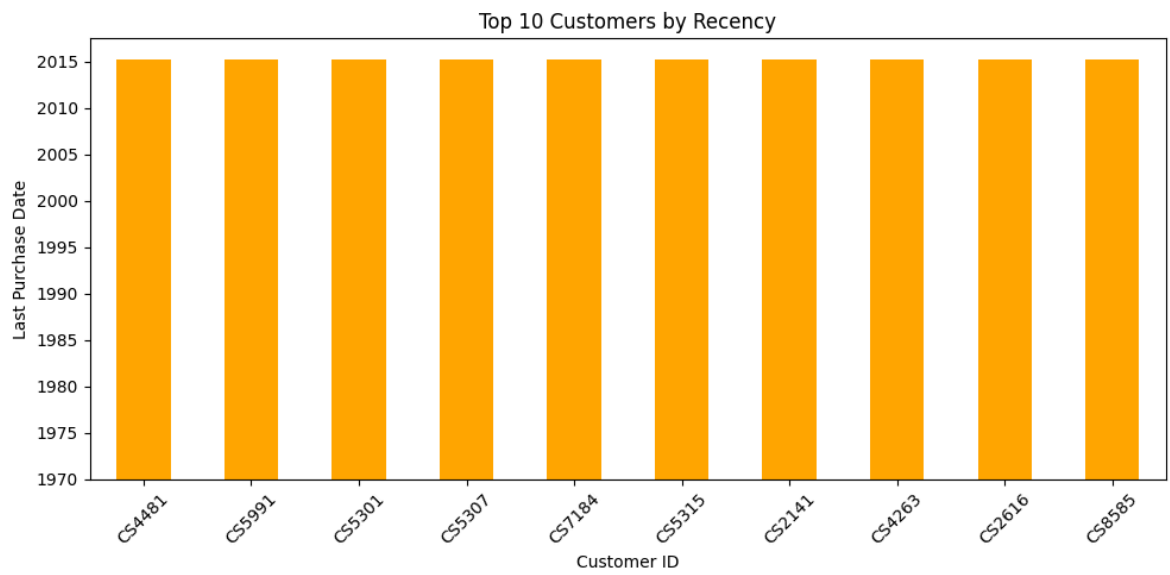
```
In [57]: top_frequent = df['customer_id'].value_counts().head(10)
```

```
plt.figure(figsize=(10, 5))
top_frequent.plot(kind='bar', color='steelblue')
plt.title("Top 10 Customers by Order Count")
plt.xlabel("Customer ID")
plt.ylabel("Number of Orders")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
In [58]: recent = df.groupby('customer_id')['trans_date'].max().sort_values(ascending=False)
```

```
plt.figure(figsize=(10, 5))
recent.index = recent.index.astype(str)
recent = recent.sort_values() # So most recent is rightmost
recent.plot(kind='bar', color='orange')
plt.title("Top 10 Customers by Recency")
plt.xlabel("Customer ID")
plt.ylabel("Last Purchase Date")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
In [63]: df.to_csv('MainData.csv')
```

```
In [65]: new.to_csv('AdditionalAnalysis.csv')
```

```
In [ ]:
```